
Churn Prediction Web Application – Comprehensive Documentation

Problem Statement

Customer churn is a critical metric for telecom companies as it directly impacts profitability. The goal of this project is to **predict customer churn** using historical customer information and service usage patterns. Accurate churn prediction allows companies to proactively engage with high-risk customers, thereby improving retention.

Data Overview

The dataset includes customer-level information such as:

- **Demographics:** Gender, SeniorCitizen, Partner, Dependents
 - **Service Info:** PhoneService, InternetService, StreamingTV, etc.
 - **Account Info:** Contract type, MonthlyCharges, Tenure, PaymentMethod
 - **Target Variable:** Churn (Yes/No)
-

Data Preprocessing Pipeline

To ensure clean and reliable input for machine learning models, a robust preprocessing pipeline was developed:

Steps:

- **Type Casting:** Convert SeniorCitizen to categorical, handle missing TotalCharges with coercion and drop NaNs.
 - **Encoding:** One-hot encoding for 16 categorical variables.
 - **Scaling:** RobustScaler used on numerical columns (tenure, MonthlyCharges, TotalCharges) to manage outliers.
 - **Consistency:** Saved feature columns and scaler are reused in prediction to maintain parity between training and inference.
-

Model Development

Multiple machine learning algorithms were evaluated using cross-validation and hyperparameter tuning:

✅ Models Built:

- Logistic Regression
- Decision Tree
- Random Forest
- Extra Trees
- K-Nearest Neighbors
- Gradient Boosting
- Support Vector Machines
- Gaussian Naive Bayes

✅ Evaluation Metrics:

- Accuracy
- AUC-ROC
- Confusion Matrix

Each model was trained on the same processed data and saved as .pkl files for later inference.

📊 Evaluation Highlights

Model	Accuracy	AUC-ROC
Logistic Regression	0.8032	0.8437
Random Forest	0.8214	0.8596
Gradient Boosting	0.8301	0.8723
Extra Trees	0.8289	0.8692
SVM (Linear)	0.7897	0.8314




✅ **Gradient Boosting** emerged as the best performer with the highest AUC.

🚀 Deployment: Streamlit Web App

A fully functional and interactive **Streamlit-based web application** was developed for end-to-end prediction.

🔍 Key Features:

- 📁 **File Upload:** Accepts raw customer data in CSV format.
- 📄 **Preview:** Displays uploaded data, schema, nulls, and descriptive stats.
- 🧠 **Model Predictions:** Automatically applies **all trained models** to predict churn.

-  **Visualizations:** AUC-ROC curves plotted dynamically for each model.
-  **Download:** Per-model prediction CSV export available.
-  **Best Model Selection:** Intelligent capture of the best-performing model.

UI Innovations:

- Modern **fade-in animations** for headings
- Tab-based architecture (Summary, All Predictions, model-specific tabs)
- Clean Plotly visuals for ROC curves
- DataFrames are styled and scrollable

Innovations & Uniqueness

Innovation	Description
Multi-model comparison	Evaluates multiple models side-by-side in real-time within the web UI.
Training-inference consistency	Ensures identical preprocessing at both stages using saved scaler and feature columns.
Auto-handling categorical levels	Robust to category mismatches or missing columns by reindex with fill.
ROC Curve per model	Dynamically generates and displays ROC curves in the app using Plotly.
Export-ready predictions	Results are downloadable per model to compare performance offline.

Project Structure

churn_app/

```





├── app.py (Streamlit frontend)
├── models/
│   ├── scaler.pkl
│   ├── train_dummy_columns.pkl
│   ├── model_rf.pkl
│   ├── model_gb.pkl
│   └── ...
├── data/
└── test.csv

```

Robustness & Security

- Error handling for missing columns, corrupt models, invalid CSV files
- Download buttons wrapped with file-safe naming conventions
- Modularized architecture for easy extension (e.g., adding more models)

Future Enhancements

-  **SHAP/LIME** explainability integration
-  **Model selection toggle** for production usage
-  **Drift detection** to alert if incoming data diverges from training distribution
-  **Model retraining trigger** based on performance

Conclusion

This churn prediction app delivers a powerful yet user-friendly interface to explore customer churn through multiple lenses — with real-time predictions, visual analytics, and flexibility to handle different models. It's designed for **business stakeholders, data scientists, and analysts** alike to make informed retention strategies.

Submission Highlights (if part of competition):

Section	Details
Team Name	Vijaykumar Radhakisan Kalaskar
Project Name	Customer Churn Predictor
Link to Web App	[e.g., Streamlit Cloud URL]
GitHub Repo	[e.g., https://github.com/yourrepo]
Demo Video	[Optional link]